

BenchPress: Benchmarks Preserving Rankings Across Multilingual Post-Training Stages

Pierre Ancey | 314520 | pierre.ancey@epfl.ch
 María Grandury | 408094 | maria.grandury@epfl.ch
 Kaustubh Ponkshe | 387628 | kaustubh.ponkshe@epfl.ch
 Raghav Singhal | 387635 | raghav.singhal@epfl.ch
 BenchPress

1 Literature Review

1.1 Benchmark reliability frameworks

Recent work shows that benchmarks differ substantially in how reliably they guide language model development. Heineman et al. (2025) introduce a signal-to-noise ratio (SNR) framework where signal measures a benchmark’s ability to separate models and noise captures score variability across checkpoints or stochastic runs. They show that high-SNR benchmarks better preserve rankings and reduce scaling-law prediction error. Complementary work studies benchmark quality through ranking consistency, discriminability, and capability-alignment deviations (Qian et al., 2026), while IRT-based methods show that many benchmark items have weak discriminative power and can often be replaced by smaller, more informative subsets (Zhou et al., 2026; Li et al., 2025).

Efficiency-oriented work reaches similar conclusions. Gupta et al. (2025) propose SMART filtering to remove easy, contaminated, or redundant examples, reducing evaluation size while preserving agreement with independent model rankings. Scaling-law studies further suggest that only some evaluations provide predictable development signal: Chen et al. (2024) forecast downstream performance from smaller models, while Schaeffer et al. (2026) study scaling laws for generative metrics such as pass@ k . Together, these papers motivate our focus on identifying which multilingual benchmarks remain locally reliable across nearby post-training checkpoints, rather than treating all benchmark averages as equally meaningful.

1.2 RLHF Methods

Reinforcement Learning from Human Feedback (RLHF) is a standard post-training paradigm for aligning language models with human preferences. Typical pipelines combine supervised fine-tuning, preference-data collection, reward modeling, and

policy optimization under a constraint to a reference model (Ouyang et al., 2022; Bai et al., 2022). This improves subjective qualities such as helpfulness, harmlessness, and instruction following, which are difficult to capture with likelihood alone.

Recent variants simplify this pipeline. Direct Preference Optimization (DPO) optimizes directly on preference pairs, avoiding explicit reward-model training and online RL (Rafailov et al., 2023). Reward-model benchmarks such as RewardBench then evaluate whether reward models reliably distinguish preferred from dispreferred responses (Lambert et al., 2024). These methods make benchmark reliability central to post-training selection: small apparent gains may reflect evaluator bias, verbosity effects, or reward-model artifacts rather than robust improvements.

1.3 RLVR Methods

Reinforcement Learning with Verifiable Rewards (RLVR) instead uses objective validators such as exact-answer checks, unit tests, symbolic solvers, or task-specific verifiers. It is therefore well suited to math, code, and structured reasoning tasks (Cobbe et al., 2021; Lewkowycz et al., 2022; Shao et al., 2024). Recent systems such as DeepSeek-Math and DeepSeek-R1 show that verifiable rewards can elicit strong reasoning behavior without relying only on human-labeled reasoning traces (Shao et al., 2024; DeepSeek-AI, 2025).

However, RLVR shifts the reliability problem from human preference labels to verifiers. Validators may be sparse, brittle, incomplete, or exploitable, and final-answer accuracy may fail to distinguish robust reasoning from lucky sampling or benchmark-specific optimization (Zheng et al., 2025; Yue et al., 2025). For RLVR, benchmark reliability should therefore measure whether verified scores provide stable, discriminative, and generalizable signal across nearby post-trained models.

References

- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, Nicholas Joseph, Saurav Kadavath, Jackson Kernion, Tom Conerly, Sheer El-Showk, Nelson Elhage, Zac Hatfield-Dodds, Danny Hernandez, Tristan Hume, Scott Johnston, Shauna Kravec, Liane Lovitt, Neel Nanda, Catherine Olsson, Chris Olah, Nicholas Schiefer, Tom Tran-Johnson, Dario Amodei, Tom Brown, Jack Clark, Sam McCandlish, Christopher Olah, Ben Mann, and Jared Kaplan. 2022. [Training a helpful and harmless assistant with reinforcement learning from human feedback](#). *arXiv preprint arXiv:2204.05862*.
- Y. Chen et al. 2024. [Scaling laws for predicting downstream performance in llms](#). *arXiv preprint arXiv:2410.08527*.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. [Training verifiers to solve math word problems](#). *arXiv preprint arXiv:2110.14168*.
- DeepSeek-AI. 2025. [Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning](#). *arXiv preprint arXiv:2501.12948*.
- Vipul Gupta, Candace Ross, David Pantoja, Rebecca J. Passonneau, Megan Ung, and Adina Williams. 2025. [Improving model evaluation using smart filtering of benchmark datasets](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics*.
- David Heineman, Valentin Hofmann, Ian Magnusson, Yuling Gu, Noah A. Smith, Hannaneh Hajishirzi, Kyle Lo, and Jesse Dodge. 2025. [Signal and noise: A framework for reducing uncertainty in language model evaluation](#). *arXiv preprint arXiv:2508.13144*.
- Nathan Lambert, Valentina Pyatkin, Jacob Morrison, Lester James V. Miranda, Bill Yuchen Lin, Khyathi Raghavi Chandu, Nouha Dziri, Sachin Kumar, Tom Zick, Yejin Choi, Noah A. Smith, and Hannaneh Hajishirzi. 2024. [Rewardbench: Evaluating reward models for language modeling](#). *arXiv preprint arXiv:2403.13787*.
- Aitor Lewkowycz, Anders Johan Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay V. Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam Neyshabur, Guy Gur-Ari, and Vedant Misra. 2022. [Solving quantitative reasoning problems with language models](#). In *Advances in Neural Information Processing Systems*.
- Peiyu Li, Xiuxiu Tang, Si Chen, Ying Cheng, Ronald Metoyer, Ting Hua, and Nitesh V. Chawla. 2025. [Adaptive testing for llm evaluation: A psychometric alternative to static benchmarks](#). *arXiv preprint arXiv:2511.04689*.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. 2022. [Training language models to follow instructions with human feedback](#). In *Advances in Neural Information Processing Systems*.
- Q. Qian et al. 2026. [Benchmark²: Systematic evaluation of llm benchmarks](#). *arXiv preprint arXiv:2601.03986*.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. 2023. [Direct preference optimization: Your language model is secretly a reward model](#). In *Advances in Neural Information Processing Systems*.
- Rylan Schaeffer, Noam Levi, Brando Miranda, and Sanmi Koyejo. 2026. [Pretraining scaling laws for generative evaluations of language models](#). In *International Conference on Learning Representations*.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. 2024. [Deepseekmath: Pushing the limits of mathematical reasoning in open language models](#). *arXiv preprint arXiv:2402.03300*.
- Yang Yue et al. 2025. Does reinforcement learning really incentivize reasoning capacity in llms beyond the base model? *arXiv preprint*.
- Chujie Zheng et al. 2025. Processbench: Identifying process errors in mathematical reasoning. *arXiv preprint*.
- Hongli Zhou, Hui Huang, Ziqing Zhao, Lvyuan Han, Huicheng Wang, Kehai Chen, Muyun Yang, Wei Bao, Jian Dong, Bing Xu, Conghui Zhu, Hailong Cao, and Tiejun Zhao. 2026. [Lost in benchmarks? rethinking large language model benchmarking with item response theory](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*.